

Compulsory exercise 2: Group 13

TMA4268 Statistical Learning V2020

Vemund Tjessem, Erik Andre Klepp Vik

23 mars, 2020

```
# install.packages('knitr') #probably already installed  
# install.packages('rmarkdown') #probably already installed  
# install.packages('ggplot2') #plotting with ggplot install.packages('ggfortify')  
# install.packages('MASS') install.packages('dplyr')  
library(knitr)  
library(rmarkdown)  
library(ggplot2)  
library(ggfortify)  
library(MASS)  
library(dplyr)  
library(ISLR)  
library(leaps)  
library(glmnet)  
library(tree)  
library(randomForest)
```

Problem 1

a)

The ridge regression coefficients β_{Ridge} are the ones that minimize

$$RSS + \lambda \sum_{j=1}^p \beta_j^2 \quad (1)$$

with $\lambda > 0$ being a tuning parameter. From before we know that

$$RSS = \sum_{i=1}^n \left(y_i - \hat{\beta}_0 - \sum_{j=1}^p \hat{\beta}_j x_{ij} \right)^2 \quad (2)$$

This can be written in terms of matrices and vectors as

$$(y - X\hat{\beta}_{Ridge})^\top (y - X\hat{\beta}_{Ridge}) + \lambda \hat{\beta}_{Ridge}^\top \hat{\beta}_{Ridge} \quad (3)$$

Differentiating this with respect to $\hat{\beta}_{Ridge}$ and setting equal to 0 gives

$$-2X^T(y - X\hat{\beta}_{Ridge}) + 2\lambda\hat{\beta}_{Ridge} = 0 \quad (4a)$$

$$X^T X \hat{\beta}_{Ridge} + \lambda \hat{\beta}_{Ridge} = X^T y \quad (4b)$$

$$\hat{\beta}_{Ridge} = (X^T X + \lambda I)^{-1} X^T y \quad (4c)$$

Where I is the identity matrix. Even though I is an identity matrix the first element would have to be 0 in order to not penalize the intercept β_0 with the ridge regression. Maybe this type of matrix has a name?

b)

Already know that the expectation value of $y = X\beta + \epsilon$ is $E[y] = X\beta$, as $E[\epsilon] = 0$. The expectation value of $\hat{\beta}_{Ridge}$ is then

$$E[\hat{\beta}_{Ridge}] = (X^T X + \lambda I)^{-1} X^T E[y] \quad (5a)$$

$$= (X^T X + \lambda I)^{-1} X^T X \beta \quad (5b)$$

This is a biased estimator as long as $\lambda \neq 0$.

The variance covariance matrix of y is $\text{Var}[y] = \text{Var}[X\beta] + \text{Var}[\epsilon] = \sigma^2$.

$$\text{Var}[\hat{\beta}_{Ridge}] = \text{Var}[(X^T X + \lambda I)^{-1} X^T y] \quad (6a)$$

$$= (X^T X + \lambda I)^{-1} X^T \text{Var}[y] [(X^T X + \lambda I)^{-1} X^T]^T \quad (6b)$$

$$= \sigma^2 (X^T X + \lambda I)^{-1} X^T X [(X^T X + \lambda I)^{-1}]^T \quad (6c)$$

c)

- (i) True
- (ii) False
- (iii) False
- (iv) True

d)

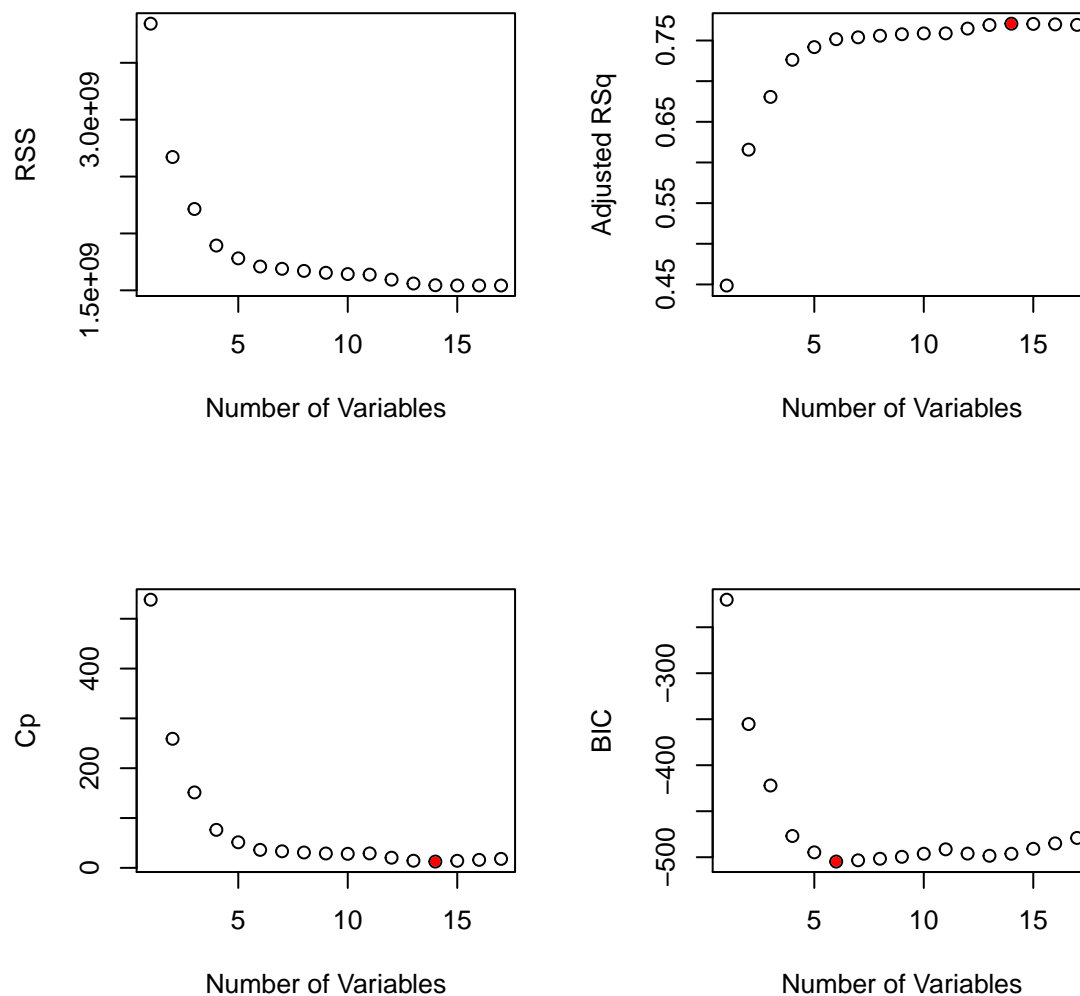
Forward selection will be performed with `Outstate` as response using the `regsubsets` function.

```
set.seed(1)
train.ind = sample(1:nrow(College), 0.5 * nrow(College))
college.train = College[train.ind, ]
college.test = College[-train.ind, ]
n_predictors = dim(College)[2] - 1
fwd.fit = regsubsets(Outstate ~ ., college.train, nvmax = n_predictors, method = "forward")
fwd.fit.summary = summary(fwd.fit)
par(mfrow = c(2, 2))
plot(fwd.fit.summary$rss, xlab = "Number of Variables", ylab = "RSS", type = )
```

```

plot(fwd.fit.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq")
fwd_best_adj2 = which.max(fwd.fit.summary$adjr2)
points(fwd_best_adj2, fwd.fit.summary$adjr2[fwd_best_adj2], col = "red", cex = 1,
       pch = 20)
plot(fwd.fit.summary$cp, xlab = "Number of Variables", ylab = "Cp")
fwd_best_cp = which.min(fwd.fit.summary$cp)
points(fwd_best_cp, fwd.fit.summary$cp[fwd_best_cp], col = "red", cex = 1, pch = 20)
fwd_best_bic = which.min(fwd.fit.summary$bic)
plot(fwd.fit.summary$bic, xlab = "Number of Variables", ylab = "BIC")
points(fwd_best_bic, fwd.fit.summary$bic[fwd_best_bic], col = "red", cex = 1, pch = 20)

```



```


predict.regsubsets = function(object, newdata, id, ...) {
  form = as.formula(object$call[[2]])
  mat = model.matrix(form, newdata)
  coefi = coef(object, id = id)
  xvars = names(coefi)

```

```

    mat[, xvars] %*% coefi
}
k = 10
set.seed(1)
folds = sample(1:k, nrow(college.train), replace = TRUE)
cv.errors = matrix(NA, k, n_predictors, dimnames = list(NULL, paste(1:n_predictors)))
# Perform CV
for (j in 1:k) {
  best_subset_method = regsubsets(Outstate ~ ., data = college.train[folds != j,
    ], nvmax = n_predictors, method = "forward")
  for (i in 1:n_predictors) {
    pred = predict(best_subset_method, college.train[folds == j, ], id = i)
    cv.errors[j, i] = mean((college.train$Outstate[folds == j] - pred)^2)
  }
}
# Compute mean cv errors for each model size
mean.cv.errors = apply(cv.errors, 2, mean)
# mean.cv.errors Plot the mean cv errors
par(mfrow = c(1, 1))
plot(mean.cv.errors, type = "b")
min_cverror = which.min(mean.cv.errors)
points(min_cverror, mean.cv.errors[min_cverror], col = "red", cex = 1, pch = 20)

```



```

# Calculating the MSE for model with 6 predictors
x.test = model.matrix(Outstate ~ ., data = college.test)
coef6 = coef(fwd.fit, id = 6)
names(coef6)

```

```
## [1] "(Intercept)" "PrivateYes" "Room.Board" "Terminal" "perc.alumni"
## [6] "Expend" "Grad.Rate"
```

```
pred = x.test[, names(coef6)] %*% coef6
MSE.forward = mean((college.test$Outstate - pred)^2)
MSE.forward
```

```
## [1] 3844857
```

The obvious choice might be the model with 14 predictors, as this had both the highest adjusted R^2 and the smallest C_p . However, since the improvement is very small for the larger models, we will choose the model with 6 variables, as this has the smallest BIC, and BIC is defined in a way that normally favors a smaller model. Cross validation also shows that 6 would be a good choice. It is not the one with the lowest mean error, but it is quite good compared to the rest. The model with 6 variables has a MSE of 3.8448572×10^6 .

e)

Model selection using the Lasso method. Since the package `glmnet` does not use the model formula language we need to set up `x` and `y`.

```
x.train = model.matrix(Outstate ~ ., data = college.train)[, -1] # -1 is to remove intercept
y.train = college.train$Outstate
x.test = model.matrix(Outstate ~ ., data = college.test)[, -1]
y.test = college.test$Outstate
lasso.fit = glmnet(x.train, y.train, alpha = 1) # alpha = 1 gives the Lasso method
set.seed(1)
lasso.fit.cv = cv.glmnet(x.train, y.train, alpha = 1)
lasso.lambda = lasso.fit.cv$lambda.1se
lasso.pred = predict(lasso.fit, s = lasso.lambda, newx = x.test)
MSE.lasso = mean(as.numeric((lasso.pred - y.test)^2))
lasso.coefs = coef(lasso.fit, s = lasso.lambda)
lasso.coefs
```

```
## 18 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) -443.5058848
## PrivateYes  1937.1756487
## Apps        .
## Accept      .
## Enroll      .
## Top10perc   13.5120398
## Top25perc   .
## F.Undergrad .
## P.Undergrad .
## Room.Board  1.0382836
## Books       .
## Personal    -0.0232975
## PhD         .
## Terminal    12.5668504
## S.F.Ratio   -46.5291891
## perc.alumni 47.4300464
## Expend      0.1486914
## Grad.Rate   26.5709586
```

```
nonzero.names = rownames(coef(lasso.fit.cv, s = lasso.lambda))[coef(lasso.fit.cv,
s = lasso.lambda)[, 1] != 0]
```

Used the function `cv.glmnet` to perform 10 fold cross validation and choose a value for λ . Instead of choosing the model with the lowest MSE in the cross validation, which used all the predictors, we chose the value `lamdba.1se` which is the largest value of λ which gives an error within 1 standard error of the minimum. The value was $\lambda = 367.77$. The reason for this is that it is a much smaller model, which only uses 8 predictors. The predictors were PrivateYes, Top10perc, Room.Board, Personal, Terminal, S.F.Ratio, perc.alumni, Expend, Grad.Rate. The MSE was 3.9033653×10^6 .

Problem 2

a)

- (i) False
- (ii) False
- (iii) True
- (iv) True (fra video om smoothing splines)

b)

The basis functions are

$$b_1(x) = x^1 \quad (7a)$$

$$b_2(x) = x^2 \quad (7b)$$

$$b_3(x) = x^3 \quad (7c)$$

$$b_4(x) = (x - q_1)_+^3 \quad (7d)$$

$$b_5(x) = (x - q_2)_+^3 \quad (7e)$$



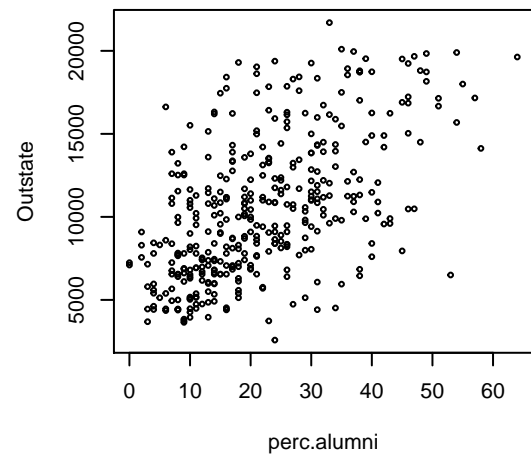
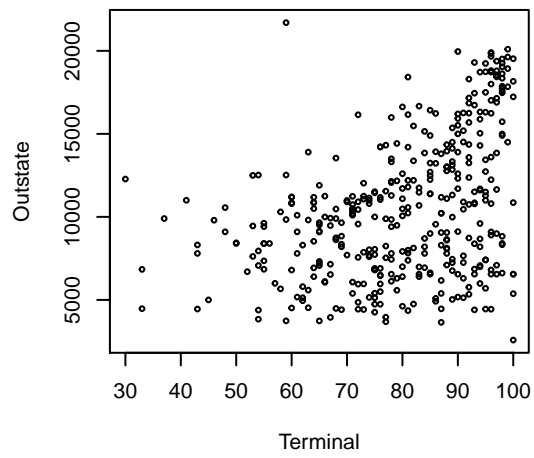


$$b_6(x) = (x - q_3)_+^3 \quad (7f)$$

c)

```
for.reg = regsubsets(Outstate ~ ., data = college.train, method = "forward")
coef.for = coef(for.reg, id = 6)
co.names = names(coef.for)[-1]
co.names[1] = "Private"
```

Will investigate the relationship between `Outstate` and the following 6 predictors: Private, Room.Board, Terminal, perc.alumni, Expend, Grad.Rate.

```
par(mfrow = c(3, 2))
plot(Outstate ~ Private, data = college.train)
plot(Outstate ~ Room.Board, data = college.train, cex = 0.5)
plot(Outstate ~ Terminal, data = college.train, cex = 0.5)
plot(Outstate ~ perc.alumni, data = college.train, cex = 0.5)
plot(Outstate ~ Expend, data = college.train, cex = 0.5)
plot(Outstate ~ Grad.Rate, data = college.train, cex = 0.5)
```




The relationship between `Outstate` and `Room.Board` seems to be approximately linear, same for `perc.alumni`. For `Terminal` on the other hand the slope seems to increase with increasing value for `Terminal`. The relation between `Outstate` and `Expend` does not seem linear, however the relation between `Outstate` and `Grad.Rate` does.

d)

- (i) Fit polynomial regression models for `Outstate` as a function of `Terminal` with polynomial degrees $d = 1, \dots, 10$.

```
par(mar = c(5.1, 4.1, 4.1, 6.5), xpd = TRUE)
degs = 10
MSE.poly.train = rep(NA, degs)
MSE.poly.test = rep(NA, degs)
plot(Outstate ~ Terminal, data = college.train, main = "Polynomial regression", cex = 0.5)
d <- seq(min(college.train$Terminal), max(college.train$Terminal), length.out = 200)
for (degree in 1:degs) {
  fm <- lm(Outstate ~ poly(Terminal, degree), data = college.train)
  assign(paste("college.train", degree, sep = "."), fm)
  lines(d, predict(fm, data.frame(Terminal = d)), col = degree)
  # Calculate training MSE
  MSE.poly.train[degree] = mean((predict(fm, college.train) - college.train$Outstate)^2)
  MSE.poly.test[degree] = mean((predict(fm, college.test) - college.test$Outstate)^2)
}
legend("topright", inset = c(-0.32, 0.2), legend = paste("d =", 1:degs), col = c(1:degs),
      lty = 1, bty = "n")
```

Polynomial regression




- (ii) Choose a suitable smoothing spline model to predict `Outstate` as a function of `Expend`.


```

x = college.train$Expend
y = college.train$Outstate
smthspl.fit = smooth.spline(x, y, cv = T)
plot(y ~ x, main = paste("Smoothing spline, df =", round(smthspl.fit$df, 3)))
lines(smthspl.fit)
points(college.test$Expend, college.test$Outstate, pch = 2, col = 2)

```

Smoothing spline, df = 4.661



```

MSE.smthspl.train = mean((predict(smthspl.fit, x)$y - y)^2)
MSE.smthspl.test = mean((predict(smthspl.fit, college.test$Expend)$y - college.test$Outstate)^2)
MSE.smthspl.train

```

```
## [1] 6871281
```


```
MSE.smthspl.test
```

```
## [1] 5943680
```

By putting `cv=T` cross validation is used to determine the degrees of freedom. They are determined to be 4.660711, higher values of `df` gives a more overfitted line.

(iii) Training MSE

```
par(mfrow = c(1, 2))
plot(1:degs, MSE.poly.train, type = "l", xlab = "Degree", ylab = "MSE train")
plot(1:degs, MSE.poly.test, type = "l", xlab = "Degree", ylab = "MSE test")
```



The training MSE for the smoothing spline is 6.8712814×10^6 .

Problem 3

a)

- (i) False
- (ii) True
- (iii) True
- (iv) False

b)

Will use random forest as there are a few strong predictors which makes the trees correlated. A random forest can help decrease the variance by decorrelating the trees. This this is a regression tree $m = p/3$ will be used. A disadvantage of using a random forest is that it does not give one tree, which makes it difficult to visualize.

```

set.seed(1)
tree.fit.randomForest = randomForest(Outstate ~ ., data = college.train, mtry = ncol(college.train)/3,
  ntree = 500, importance = TRUE)
yhat.randomForest = predict(tree.fit.randomForest, newdata = college.test)
MSE.randomForest = mean((yhat.randomForest - college.test$Outstate)^2)
importance(tree.fit.randomForest)

```

```

##           %IncMSE IncNodePurity
## Private      23.884645      572396197
## Apps         12.660213      119525827
## Accept       12.809286      113271754
## Enroll        9.960338      105290813
## Top10perc    12.597062      367734616
## Top25perc     7.042636      161501521
## F.Undergrad  11.136443      122819717
## P.Undergrad  11.097160      164097235
## Room.Board   28.209705     1012135468
## Books         2.912269       71226082
## Personal      6.834351      119306116
## PhD          9.369483      225362542
## Terminal     9.665515      196862503
## S.F.Ratio    12.615659      390623584
## perc.alumni  15.097476      337449297
## Expend       34.432054     2218648378
## Grad.Rate    15.356554      545391655


```

```

varImpPlot(tree.fit.randomForest, type = 1)

```

tree.fit.randomForest



c)

Compare MSEs of the different methods.

```
MSE.forward
```

```
## [1] 3844857
```

```
MSE.lasso
```

```
## [1] 3903365
```

```
MSE.poly.test
```

```
## [1] 11892400 11103379 10937428 10936549 11005681 11105882 16429293 10914136
```

```
## [9] 81326931 24310822
```

```
MSE.smthspl.test
```

```
## [1] 5943680
```

```
MSE.randomForest
```

```
## [1] 2566908
```